
LATEST ARTIFICIAL INTELLIGENCE PAPER STUDY

EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks

INDEX

1. COMPOUND SCALING

- WHAT IS MODEL SCALING IN DEEP LEARNING?
- INTRODUCE ABOUT COMPOUND SCALING
- THE RESULT USING COMPOUND SCALIING

2. EFFICIENTNET

- WHAT IS EFFICIENTNET?
- EFFICIENTNET'S ARCHITECTURE
- COMPARE WITH OTHER MODELS



1. COMPOUND SCALING

WHAT IS MODEL SCALING IN DEEP LEARNING?

- **Model scaling is that changes model's width, depth, resolution to improve model's performance.**
 - **Width: the number of channels**
 - **Depth: the number of layers**
 - **Resolution: up pixel numbers of input image**
-

THE EFFECT OF SCALING

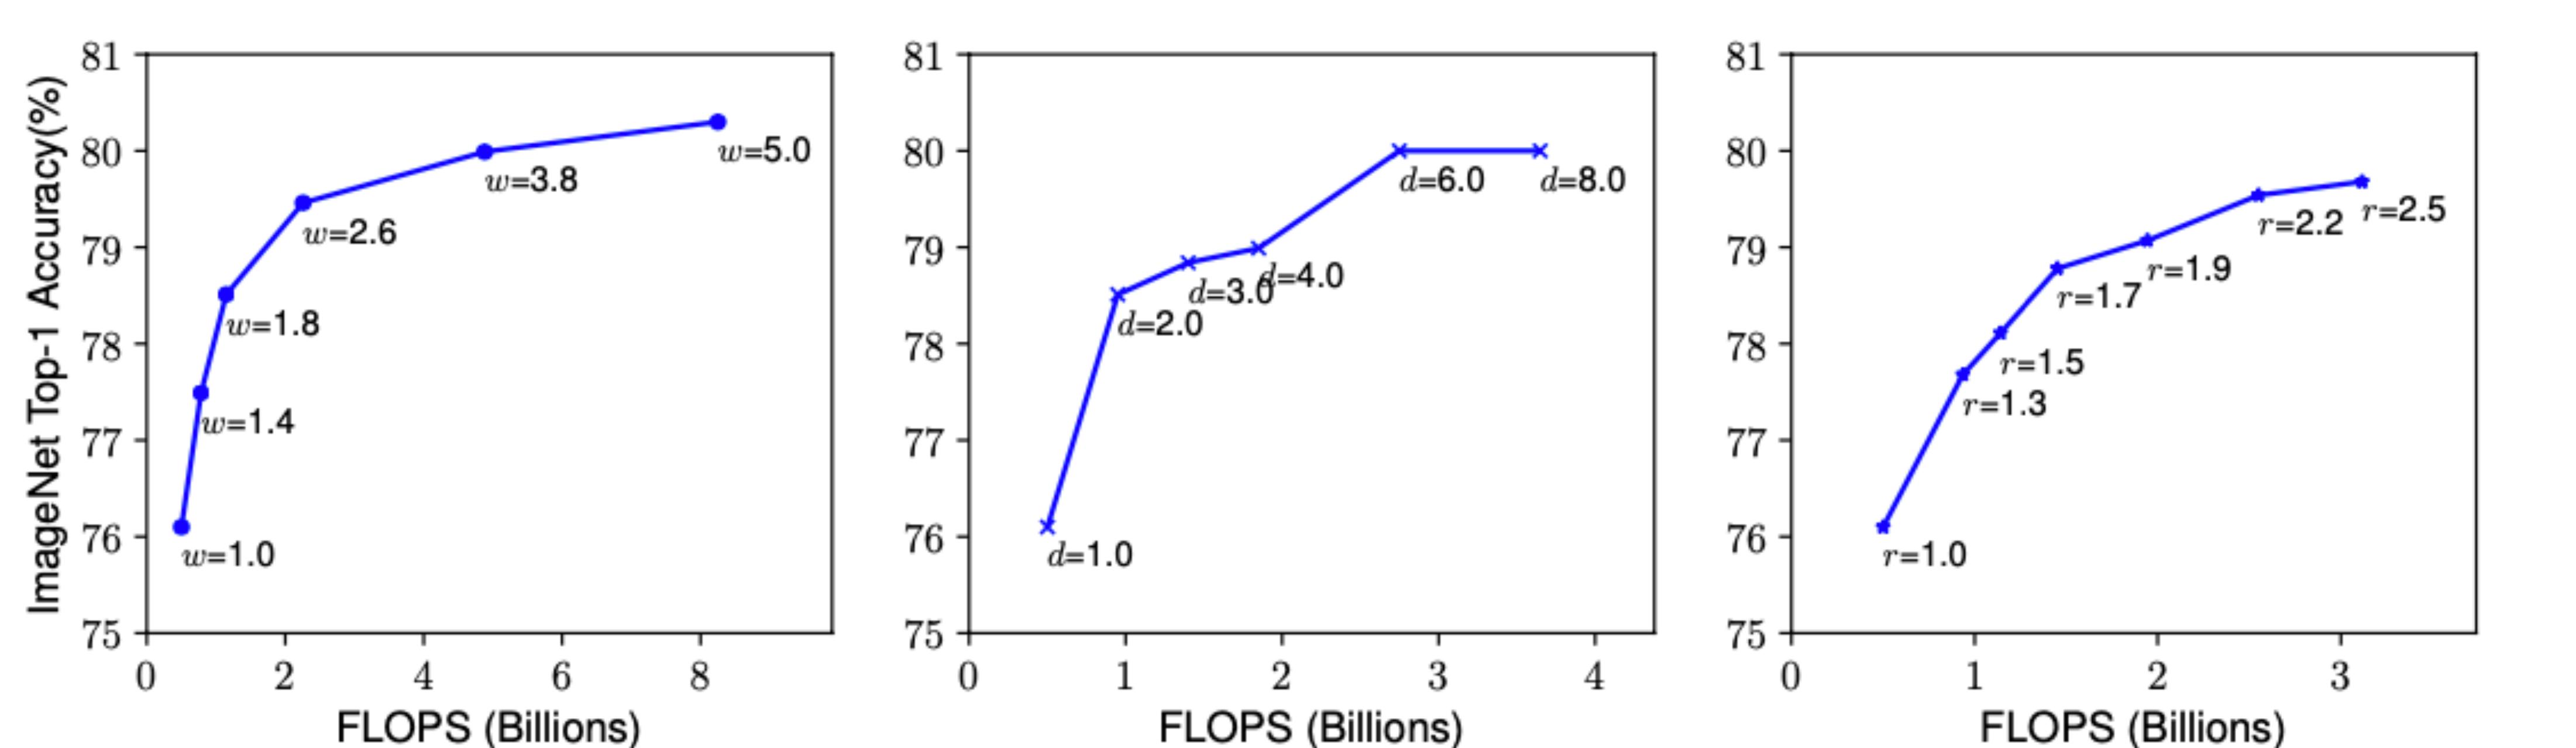
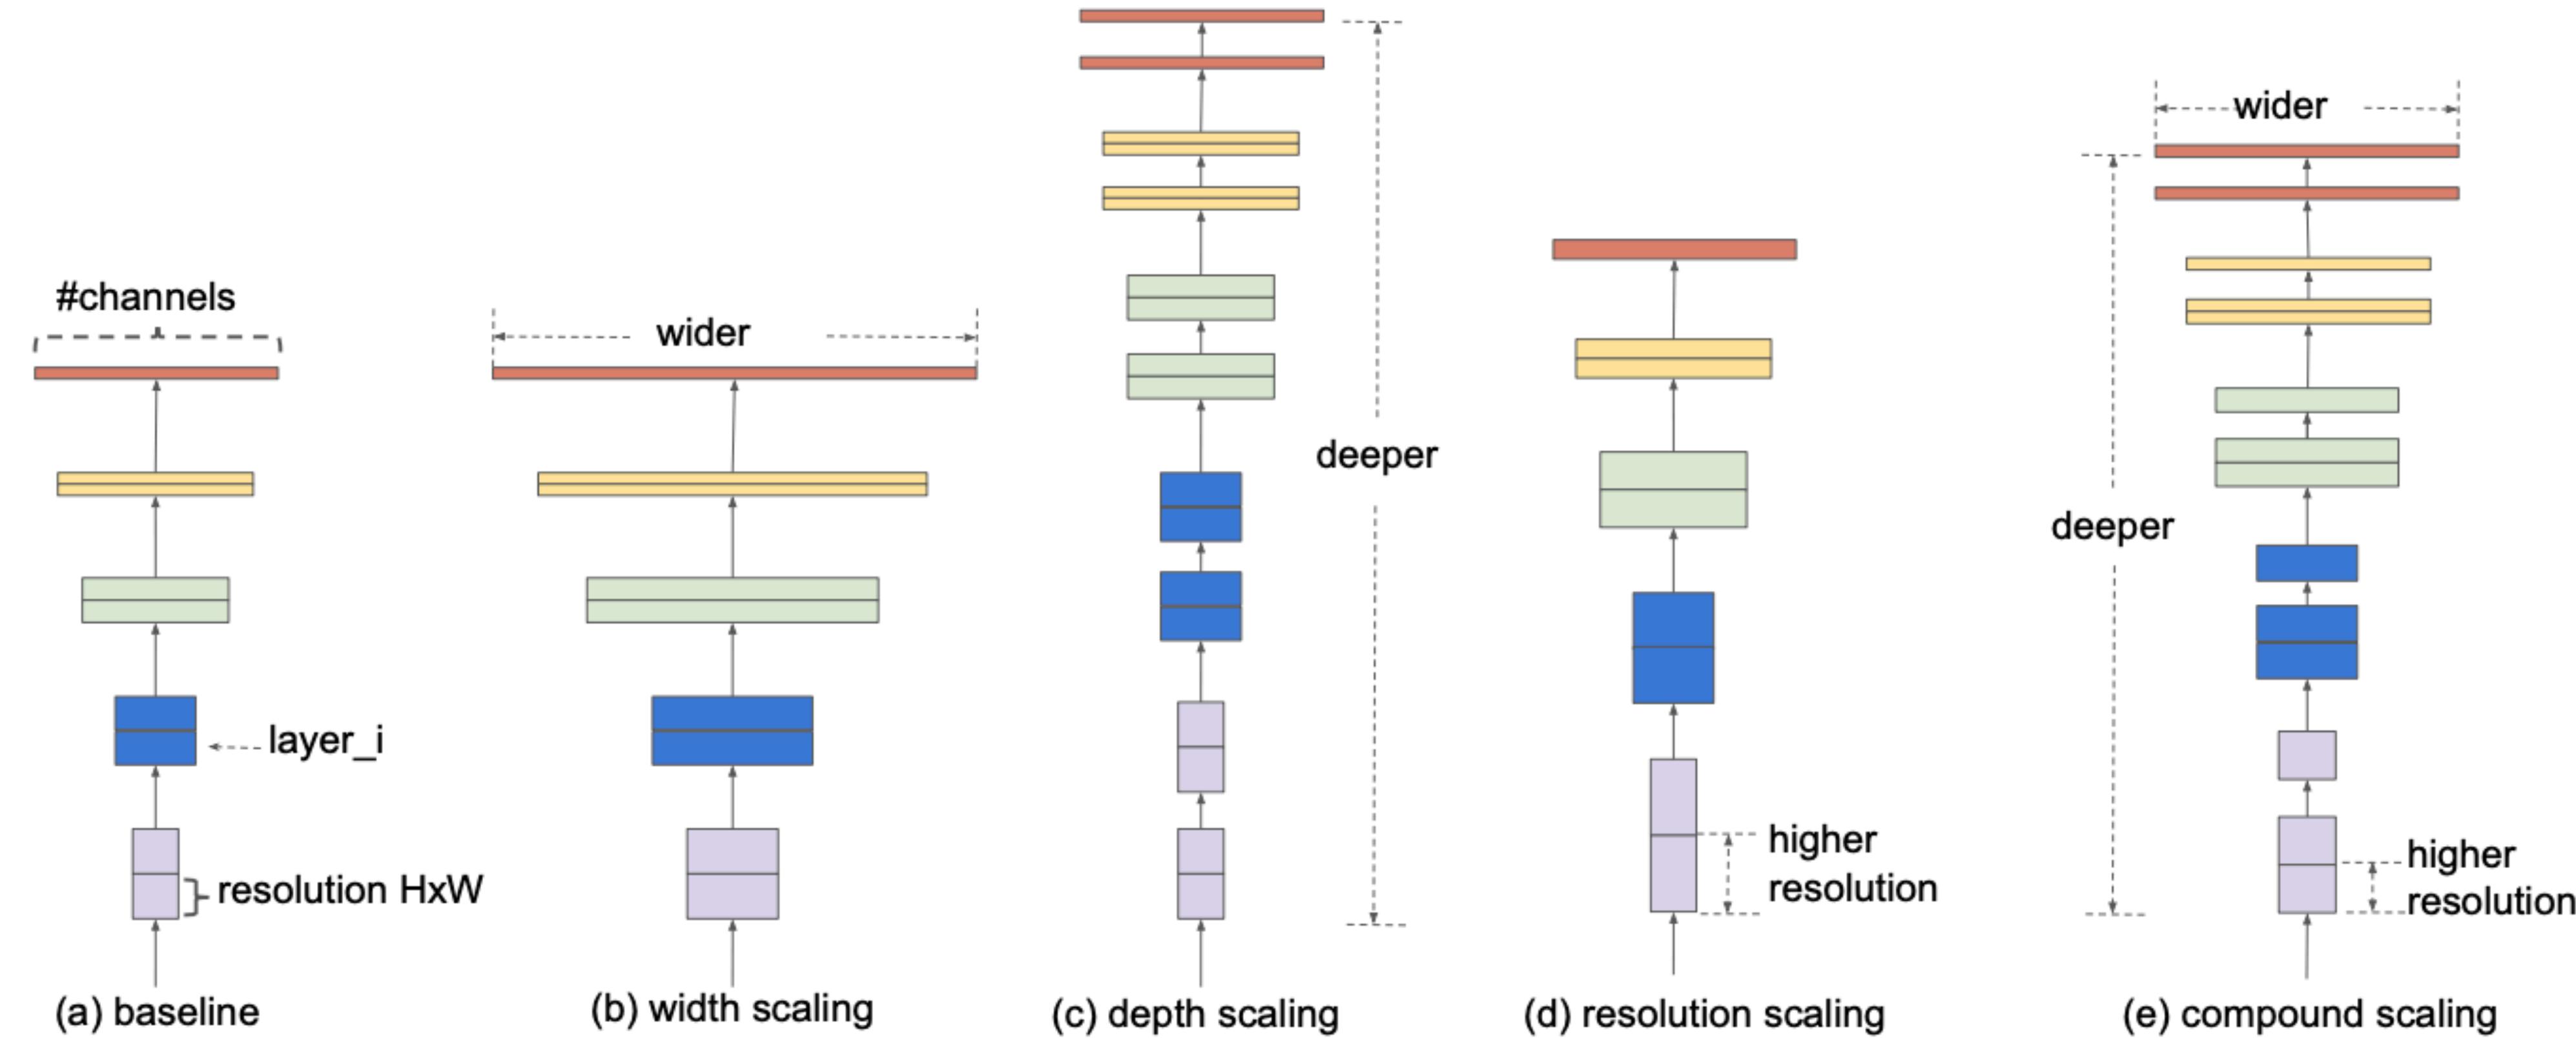


Figure 3. Scaling Up a Baseline Model with Different Network Width (w), Depth (d), and Resolution (r) Coefficients. Bigger networks with larger width, depth, or resolution tend to achieve higher accuracy, but the accuracy gain quickly saturates after reaching 80%, demonstrating the limitation of single dimension scaling. Baseline network is described in Table 1.

FLOPS : floating point operations per second

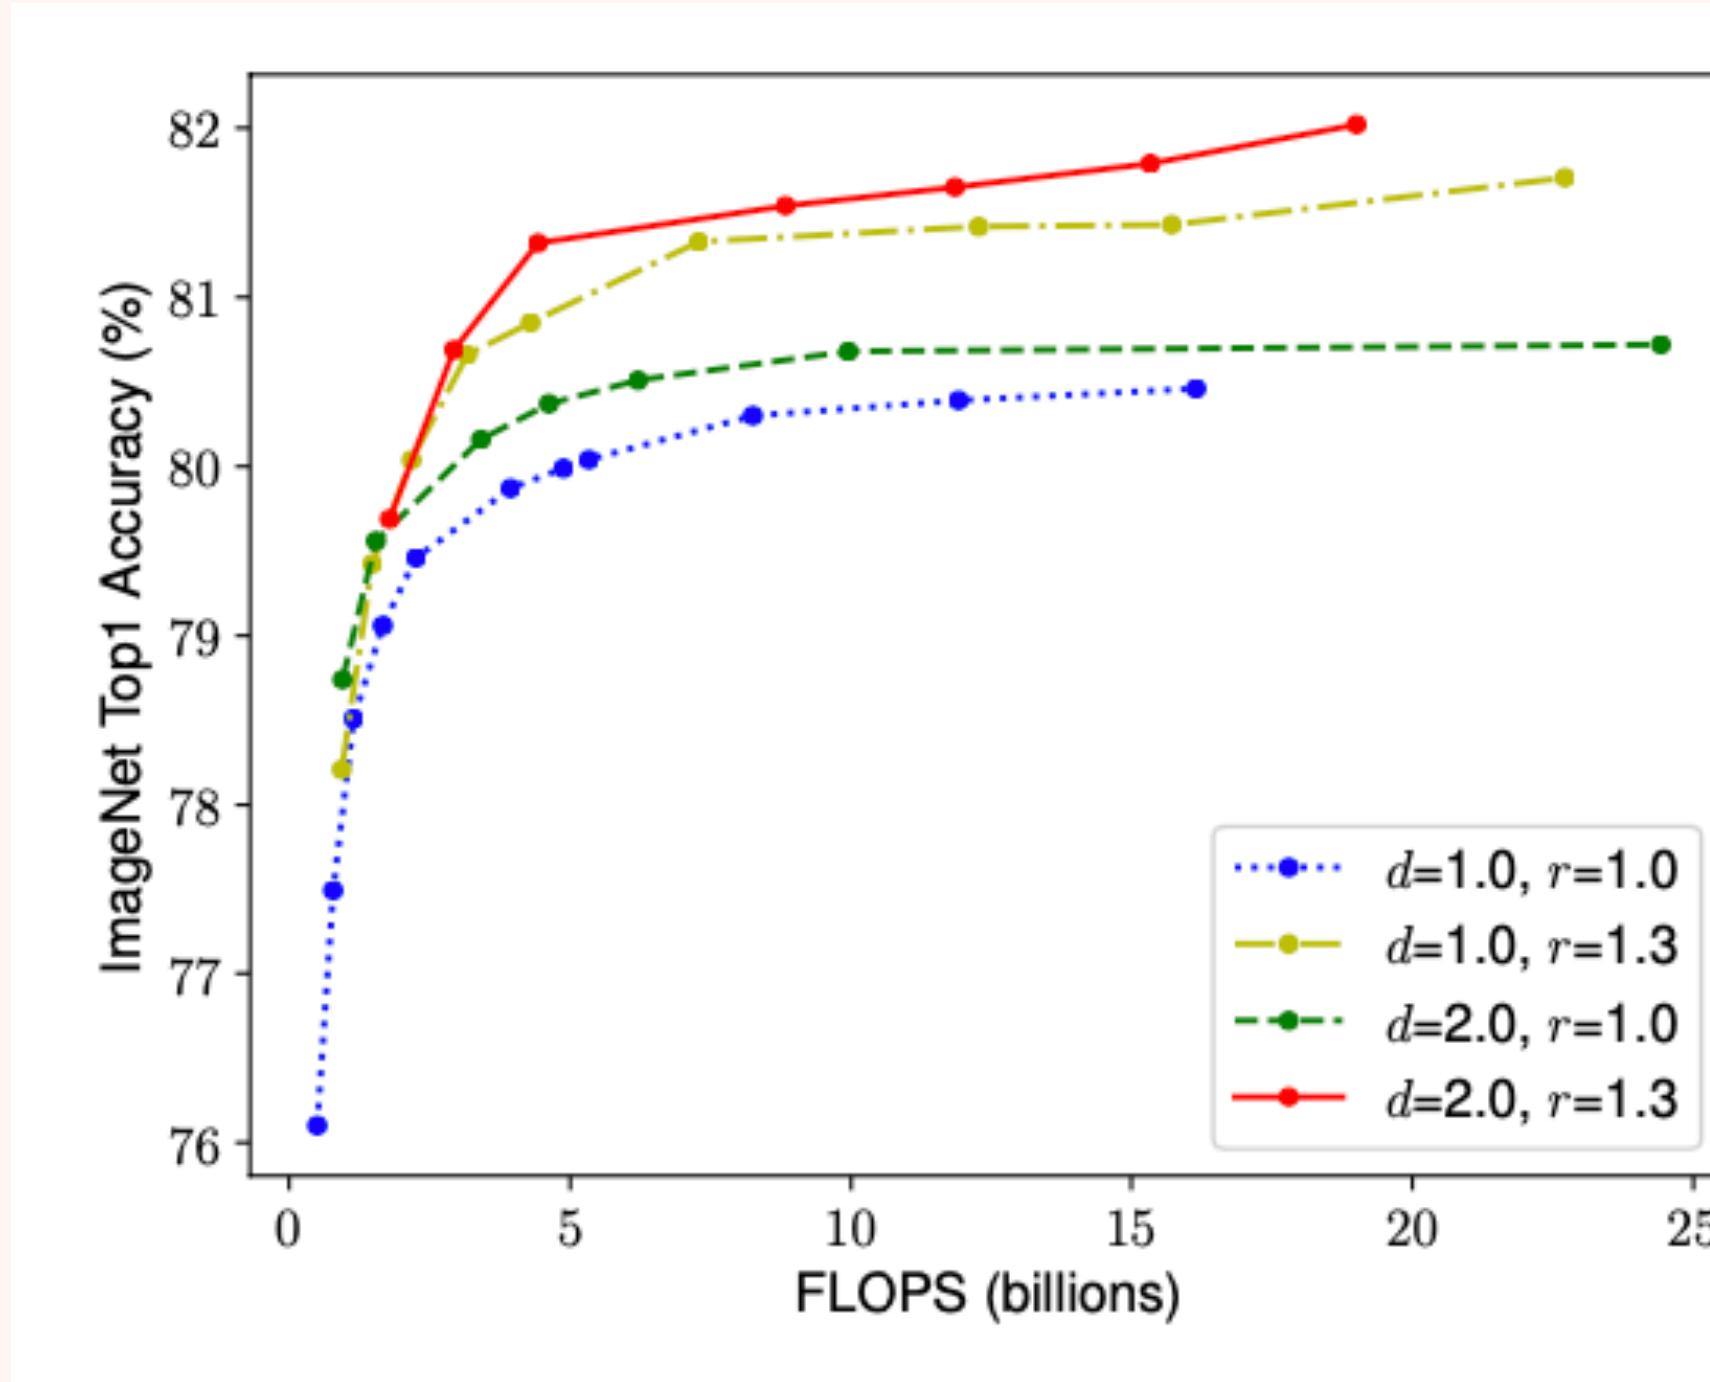
HOW TO FIND PERFECT BALANCE?



THE ANSWER IS COMPOUND SCALING

INTRODUCE ABOUT COMPOUND SCALING

➤ Compound scaling is a method to find good balanced scaling factors



- width is fixed
- depth same + resolution different : acc increases about 1%
- resolution same + depth different : acc increases about 0.3%
- The best fator's ratio is red line

INTRODUCE ABOUT COMPOUND SCALING

depth: $d = \alpha^\phi$

width: $w = \beta^\phi$

resolution: $r = \gamma^\phi$

s.t. $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$

$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$

➤ Compound scaling notation

Compound coefficient: ϕ (fai)

α, β, γ are constants (small grid search)

➤ 1. Select model

2. Fix $\phi = 1$

3. Find α, β, γ by small grid search

4. Fix α, β, γ

5. Control ϕ

THE RESULT USING COMPOUND SCALING

Table 3. Scaling Up MobileNets and ResNet.

Model	FLOPS	Top-1 Acc.
Baseline MobileNetV1 (Howard et al., 2017)	0.6B	70.6%
Scale MobileNetV1 by width ($w=2$)	2.2B	74.2%
Scale MobileNetV1 by resolution ($r=2$)	2.2B	72.7%
compound scale ($d=1.4$, $w=1.2$, $r=1.3$)	2.3B	75.6%
Baseline MobileNetV2 (Sandler et al., 2018)	0.3B	72.0%
Scale MobileNetV2 by depth ($d=4$)	1.2B	76.8%
Scale MobileNetV2 by width ($w=2$)	1.1B	76.4%
Scale MobileNetV2 by resolution ($r=2$)	1.2B	74.8%
MobileNetV2 compound scale	1.3B	77.4%
Baseline ResNet-50 (He et al., 2016)	4.1B	76.0%
Scale ResNet-50 by depth ($d=4$)	16.2B	78.1%
Scale ResNet-50 by width ($w=2$)	14.7B	77.7%
Scale ResNet-50 by resolution ($r=2$)	16.4B	77.5%
ResNet-50 compound scale	16.7B	78.8%

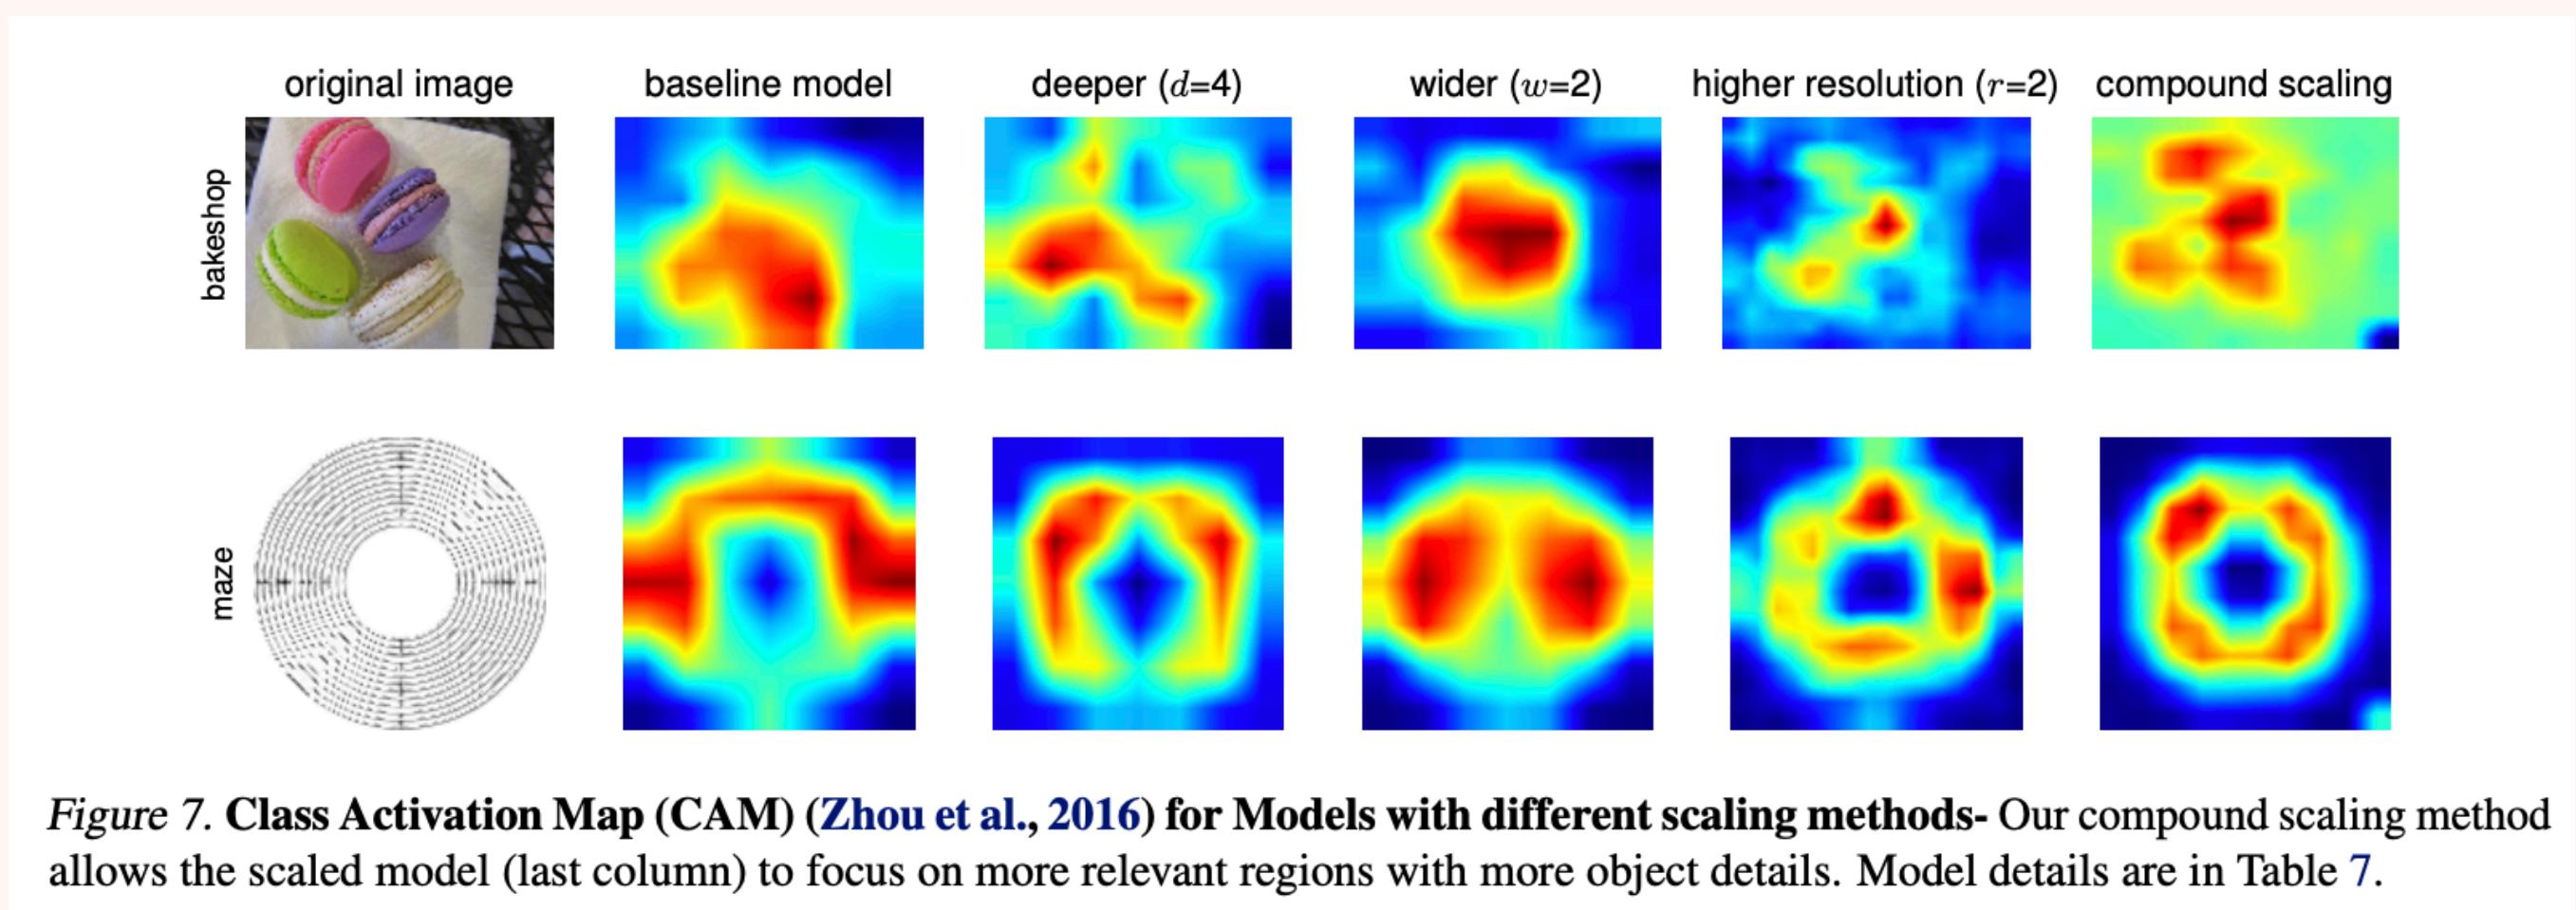


Figure 7. Class Activation Map (CAM) (Zhou et al., 2016) for Models with different scaling methods- Our compound scaling method allows the scaled model (last column) to focus on more relevant regions with more object details. Model details are in Table 7.

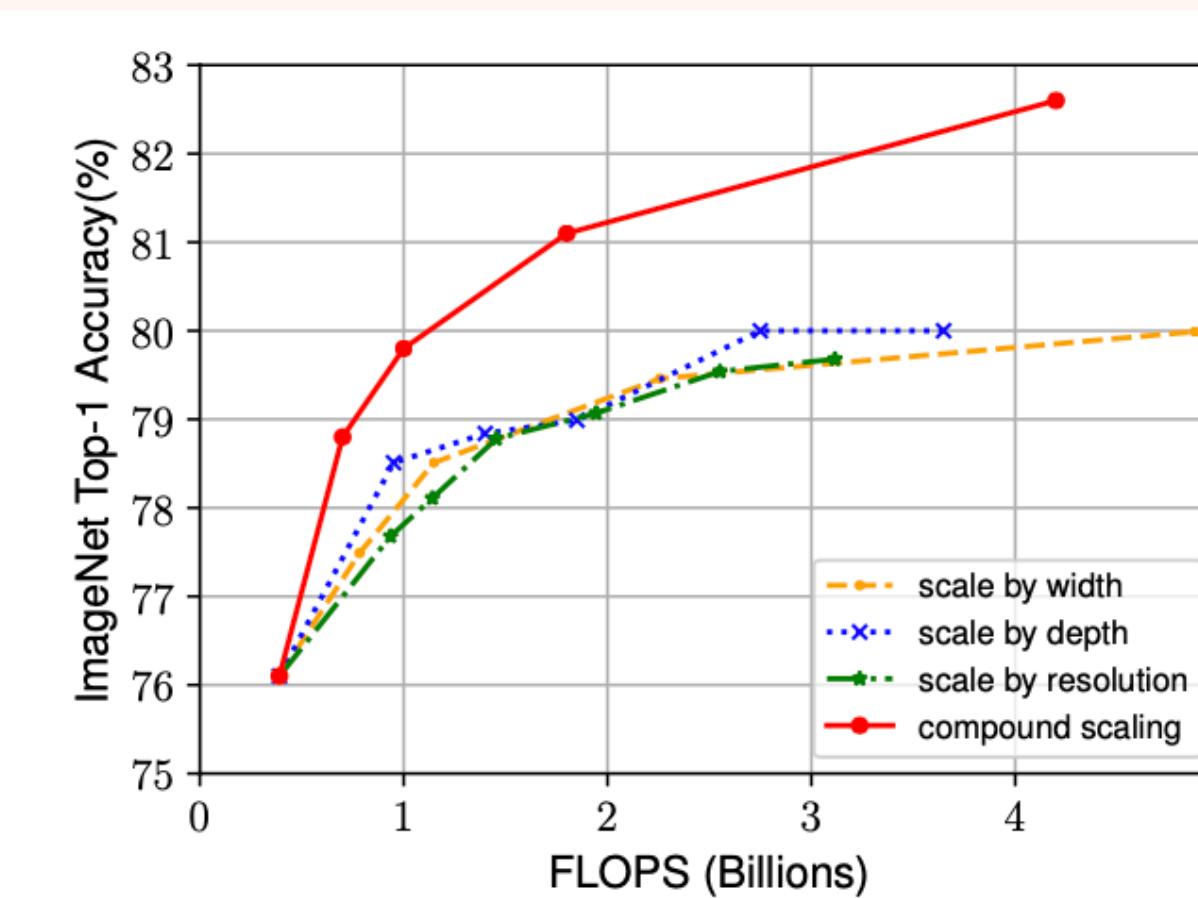


Figure 8. Scaling Up EfficientNet-B0 with Different Methods.

Table 7. Scaled Models Used in Figure 7.

Model	FLOPS	Top-1 Acc.
Baseline model (EfficientNet-B0)	0.4B	77.3%
Scale model by depth ($d=4$)	1.8B	79.0%
Scale model by width ($w=2$)	1.8B	78.9%
Scale model by resolution ($r=2$)	1.9B	79.1%
Compound Scale ($d=1.4$, $w=1.2$, $r=1.3$)	1.8B	81.1%

2. EFFICIENTNET

WHAT IS EFFICIENTNET?

- **proposed mobile-size baseline network model**
- **The model is designed for mobile from AutoML**
- **Performance good & Param's number & FLOPS down**
- **It is similar to MnasNET (neural architecture search)**

Flow: mobileNetV1 -> mobileNetV2 -> MnasNet -> efficientNet

EFFICIENTNET'S ARCHITECTURE

Table 1. EfficientNet-B0 baseline network – Each row describes a stage i with \hat{L}_i layers, with input resolution $\langle \hat{H}_i, \hat{W}_i \rangle$ and output channels \hat{C}_i . Notations are adopted from equation 2.

Stage i	Operator $\hat{\mathcal{F}}_i$	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels \hat{C}_i	#Layers \hat{L}_i
1	Conv3x3	224×224	32	1
2	MBConv1, k3x3	112×112	16	1
3	MBConv6, k3x3	112×112	24	2
4	MBConv6, k5x5	56×56	40	2
5	MBConv6, k3x3	28×28	80	3
6	MBConv6, k5x5	14×14	112	3
7	MBConv6, k5x5	14×14	192	4
8	MBConv6, k3x3	7×7	320	1
9	Conv1x1 & Pooling & FC	7×7	1280	1

EN B0	kernel size	Repeat	input filter	output filter	expand ratio	se_ratio	strides
MBConv1 (Block1)	3	1	32	16	1	0.25	1
MBConv6 (Block2)	3	2	16	24	6	0.25	2
MBConv6 (Block3)	5	2	24	40	6	0.25	2
MBConv6 (Block4)	3	3	40	80	6	0.25	2
MBConv6 (Block5)	5	3	80	112	6	0.25	1
MBConv6 (Block6)	5	4	112	192	6	0.25	2
MBConv6 (Block7)	3	1	192	320	6	0.25	1

- **Swish activation**
- **Squeeze & Excitation Block (SE Block, SE Net)**
- **Mobile Inverted Residual Block (MBConv)**

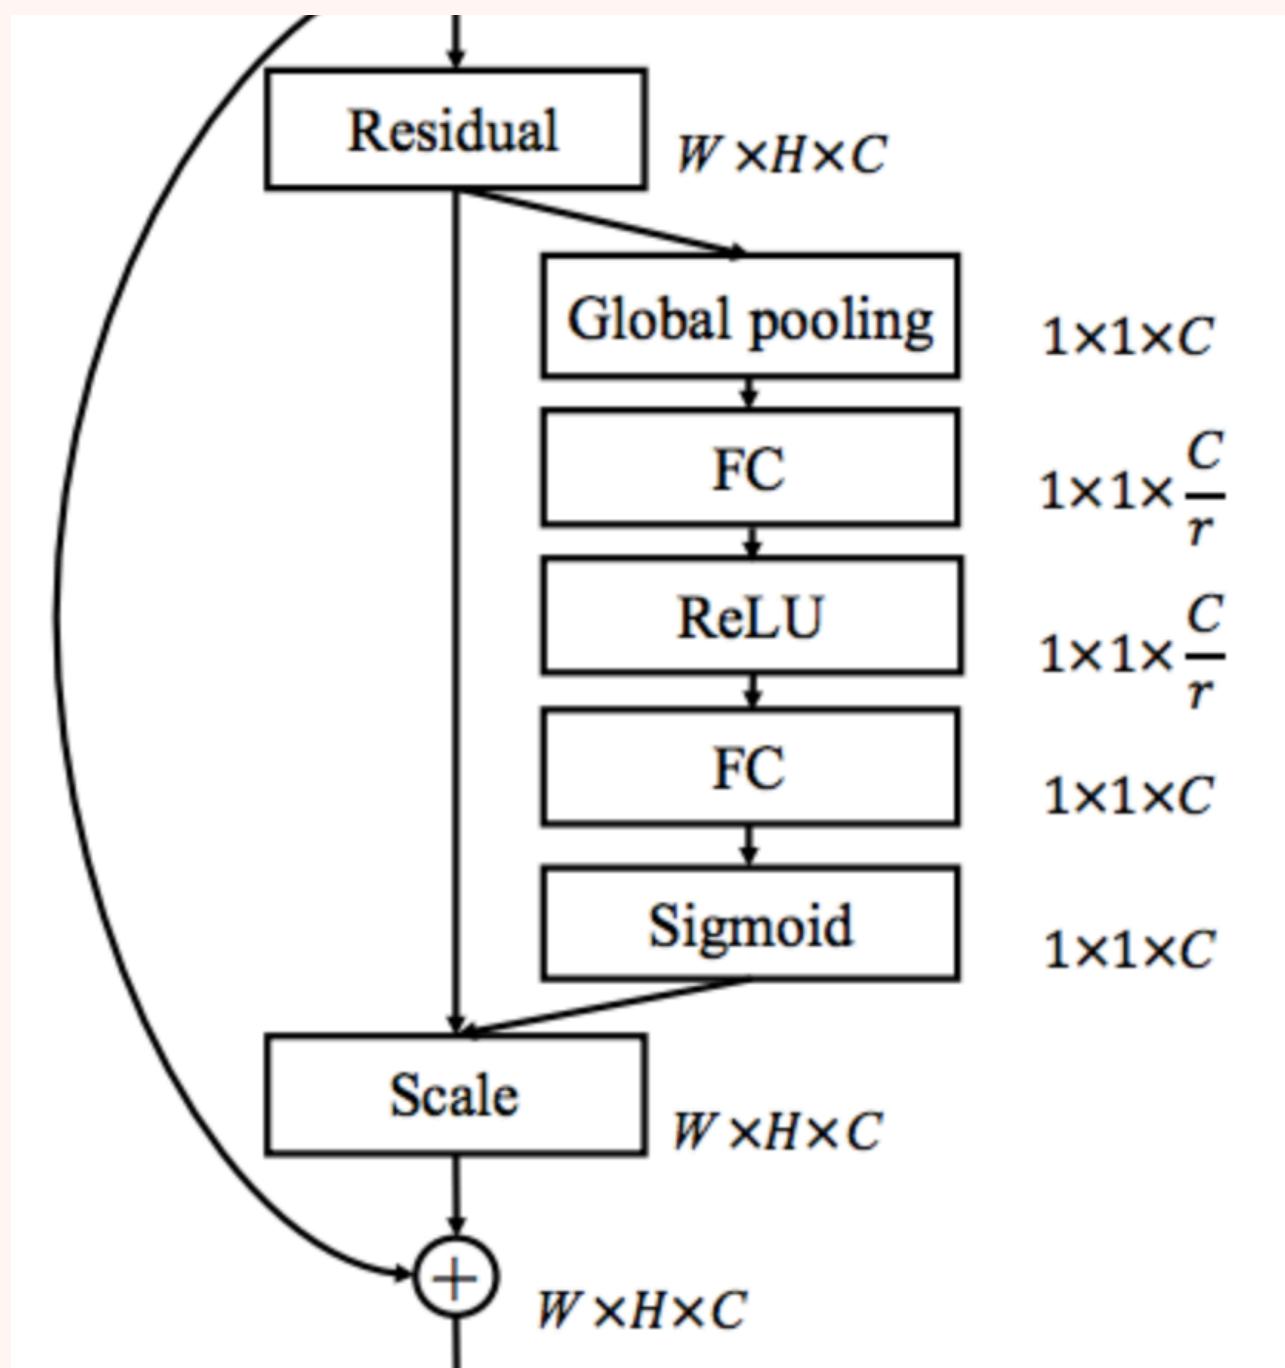
SWISH ACTIVATION

- Google brain team suggested a new activation
- It is better than Relu (solved nullifies negetive values)
- Formula : $\text{swish}(x) = x * \text{sigmoid}(x)$

```
# x * sigmoid(x)

def swish_activation(x):
    return x * (1 / (1+np.exp(-x)))
```

SQUEEZE AND EXCITATION BLOCK



- Find Each channel feature -> 1×1 feature map
- Improve Performance > increase FLOPS = useful

```
def se_block(x, filter_num, squeeze_ratio):  
    se = tf.keras.layers.GlobalAveragePooling2D()(x)  
    se = tf.keras.layers.Reshape((1,1,filter_num))(se)  
    squeezed_filters = max(1, int(filter_num * squeeze_ratio))  
    se = tf.keras.layers.Conv2D(squeezed_filters,(1,1),activation='relu')(se)  
    se = tf.keras.layers.Conv2D(filter_num,(1,1),activation='sigmoid')(se)  
    return multiply()([x, se])
```

INVERTED RESIDUAL BLOCK

```
def residual_block(x, squeeze, expand):
    re = tf.keras.layers.Conv2D(squeeze, (1,1),padding='same', activation='relu')(x)
    re = tf.keras.layers.Conv2D(squeeze, (3,3),padding='same', activation='relu')(re)
    re = tf.keras.layers.Conv2D(expand, (1,1),padding='same', activation='relu')(re)
    return Add()([re, x])
```

```
def inverted_residual_block(x, expand, squeeze):
    block = tf.keras.layers.Conv2D(expand, (1,1),padding='same',activation='relu')(x)
    block = tf.keras.layers.DepthwiseConv2D((3,3),padding='same', activation='relu')(block)
    block = tf.keras.layers.Conv2D(squeeze, (1,1),padding='same', activation='relu')(block)
    return Add()([block, x])
```

MOBILE INVERTED RESIDUAL BLOCK

```
def mbConv_block(
    input_data,repeat_num, kernel_size,input_filter,output_filter,expand_ratio,se_ratio,strides, drop_ratio):

    # expansion phase (1,1)
    expanded_filter = input_filter * expand_ratio
    x = tf.keras.layers.Conv2D(expanded_filter, 1, padding='same', use_bias=False)(input_data)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.Activation('swish')(x)

    # Depthwise convolution phase (k,k)
    x = tf.keras.layers.DepthwiseConv2D(kernel_size, strides, padding='same', use_bias=False)(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.Activation('swish')(x)

    # Squeeze and excitation phase (1,1)
    squeezed_filter = max(1, int(input_filter * se_ratio))
    se = tf.keras.layers.GlobalAveragePooling2D()(x)
    se = tf.keras.layers.Reshape((1, 1, expanded_filter))(se)
    se = tf.keras.layers.Conv2D(squeezed_filter,1)(se)
    se = tf.keras.layers.Activation('swish')(se)
    se = tf.keras.layers.Conv2D(expanded_filter,1, activation='sigmoid')(se)
    x = tf.keras.layers.multiply([x, se])

    # Output phase (1,1)
    x = tf.keras.layers.Conv2D(output_filter, 1, padding='same', use_bias=False)(x)
    x = tf.keras.layers.BatchNormalization()(x)

    if repeat_num == 1:
        pass
    else:
        x = tf.keras.layers.Dropout(drop_ratio)(x)

    return x
```

EFFICIENTNET B0'S ARCHITECTURE

Table 1. EfficientNet-B0 baseline network – Each row describes a stage i with \hat{L}_i layers, with input resolution $\langle \hat{H}_i, \hat{W}_i \rangle$ and output channels \hat{C}_i . Notations are adopted from equation 2.

Stage i	Operator $\hat{\mathcal{F}}_i$	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels \hat{C}_i	#Layers \hat{L}_i
1	Conv3x3	224×224	32	1
2	MBConv1, k3x3	112×112	16	1
3	MBConv6, k3x3	112×112	24	2
4	MBConv6, k5x5	56×56	40	2
5	MBConv6, k3x3	28×28	80	3
6	MBConv6, k5x5	14×14	112	3
7	MBConv6, k5x5	14×14	192	4
8	MBConv6, k3x3	7×7	320	1
9	Conv1x1 & Pooling & FC	7×7	1280	1

```

def Model():

    model_input = tf.keras.layers.Input(shape = (224,224,3))

    # stem
    x = tf.keras.layers.Conv2D(32, (3,3), padding='same',strides = (2,2))(model_input)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.Activation('swish')(x)

    # mbConv_blocks
    MBConv1_1 = mbConv_block(x,3,1,32,16,1,0.25,1,0.2)

    MBConv6_2_1 = mbConv_block(MBConv1_1,3,1,16,24,6,0.25,2,0.2)
    MBConv6_2_2 = mbConv_block(MBConv1_1,3,2,16,24,6,0.25,2,0.2)
    x = tf.keras.layers.add([MBConv6_2_1, MBConv6_2_2])

    MBConv6_3_1 = mbConv_block(x,5,1,24,40,6,0.25,2,0.2)
    MBConv6_3_2 = mbConv_block(x,5,2,24,40,6,0.25,2,0.2)
    x = tf.keras.layers.add([MBConv6_3_1, MBConv6_3_2])

    MBConv6_4_1 = mbConv_block(x,3,1,40,80,6,0.25,2,0.2)
    MBConv6_4_2 = mbConv_block(x,3,2,40,80,6,0.25,2,0.2)
    MBConv6_4_3 = mbConv_block(x,3,3,40,80,6,0.25,2,0.2)
    x = tf.keras.layers.add([MBConv6_4_1, MBConv6_4_2, MBConv6_4_3])

    MBConv6_5_1 = mbConv_block(x,5,1,80,112,6,0.25,1,0.2)
    MBConv6_5_2 = mbConv_block(x,5,2,80,112,6,0.25,1,0.2)
    MBConv6_5_3 = mbConv_block(x,5,3,80,112,6,0.25,1,0.2)
    x = tf.keras.layers.add([MBConv6_5_1, MBConv6_5_2, MBConv6_5_3])

    MBConv6_6_1 = mbConv_block(x,5,1,112,192,6,0.25,2,0.2)
    MBConv6_6_2 = mbConv_block(x,5,2,112,192,6,0.25,2,0.2)
    MBConv6_6_3 = mbConv_block(x,5,3,112,192,6,0.25,2,0.2)
    MBConv6_6_4 = mbConv_block(x,5,4,112,192,6,0.25,2,0.2)
    x = tf.keras.layers.add([MBConv6_6_1, MBConv6_6_2, MBConv6_6_3, MBConv6_6_4])

    MBConv6_7 = mbConv_block(x,3,1,192,320,6,0.25,1,0.2)

    # output
    output = tf.keras.layers.GlobalAveragePooling2D()(x)

    model = tf.keras.Model(model_input, output)

    return model

```

COMPARE WITH OTHER MODELS

➤ Compound scaling ratio

Model	width	depth	Resolution	dropout
EN B0	1.0	1.0	224	0.2
EN B1	1.0	1.1	240	0.2
EN B2	1.1	1.2	260	0.3
EN B3	1.2	1.4	300	0.3
EN B4	1.4	1.8	380	0.4
EN B5	1.6	2.2	456	0.4
EN B6	1.8	2.6	528	0.5
EN B7	2.0	3.1	600	0.5

Table 2. EfficientNet Performance Results on ImageNet (Russakovsky et al., 2015). All EfficientNet models are scaled from our baseline EfficientNet-B0 using different compound coefficient ϕ in Equation 3. ConvNets with similar top-1/top-5 accuracy are grouped together for efficiency comparison. Our scaled EfficientNet models consistently reduce parameters and FLOPS by an order of magnitude (up to 8.4x parameter reduction and up to 16x FLOPS reduction) than existing ConvNets.

Model	Top-1 Acc.	Top-5 Acc.	#Params	Ratio-to-EfficientNet	#FLOPS	Ratio-to-EfficientNet
EfficientNet-B0	77.3%	93.5%	5.3M	1x	0.39B	1x
ResNet-50 (He et al., 2016)	76.0%	93.0%	26M	4.9x	4.1B	11x
DenseNet-169 (Huang et al., 2017)	76.2%	93.2%	14M	2.6x	3.5B	8.9x
EfficientNet-B1	79.2%	94.5%	7.8M	1x	0.70B	1x
ResNet-152 (He et al., 2016)	77.8%	93.8%	60M	7.6x	11B	16x
DenseNet-264 (Huang et al., 2017)	77.9%	93.9%	34M	4.3x	6.0B	8.6x
Inception-v3 (Szegedy et al., 2016)	78.8%	94.4%	24M	3.0x	5.7B	8.1x
Xception (Chollet, 2017)	79.0%	94.5%	23M	3.0x	8.4B	12x
EfficientNet-B2	80.3%	95.0%	9.2M	1x	1.0B	1x
Inception-v4 (Szegedy et al., 2017)	80.0%	95.0%	48M	5.2x	13B	13x
Inception-resnet-v2 (Szegedy et al., 2017)	80.1%	95.1%	56M	6.1x	13B	13x
EfficientNet-B3	81.7%	95.6%	12M	1x	1.8B	1x
ResNeXt-101 (Xie et al., 2017)	80.9%	95.6%	84M	7.0x	32B	18x
PolyNet (Zhang et al., 2017)	81.3%	95.8%	92M	7.7x	35B	19x
EfficientNet-B4	83.0%	96.3%	19M	1x	4.2B	1x
SENet (Hu et al., 2018)	82.7%	96.2%	146M	7.7x	42B	10x
NASNet-A (Zoph et al., 2018)	82.7%	96.2%	89M	4.7x	24B	5.7x
AmoebaNet-A (Real et al., 2019)	82.8%	96.1%	87M	4.6x	23B	5.5x
PNASNet (Liu et al., 2018)	82.9%	96.2%	86M	4.5x	23B	6.0x
EfficientNet-B5	83.7%	96.7%	30M	1x	9.9B	1x
AmoebaNet-C (Cubuk et al., 2019)	83.5%	96.5%	155M	5.2x	41B	4.1x
EfficientNet-B6	84.2%	96.8%	43M	1x	19B	1x
EfficientNet-B7	84.4%	97.1%	66M	1x	37B	1x
GPipe (Huang et al., 2018)	84.3%	97.0%	557M	8.4x	-	-

We omit ensemble and multi-crop models (Hu et al., 2018), or models pretrained on 3.5B Instagram images (Mahajan et al., 2018).

THANK YOU :)
